

Executive Summary

Evolve is a SaaS educational platform designed to revolutionize first-year university learning through AI-driven personalization and project-based pedagogy. This report presents a detailed analysis of its core features, evaluates their technical viability, surveys best-in-class implementations, and recommends an optimal technology stack and development methodology.

1. Product Overview

- **Name:** Evolve
- **Objective:** Help first-year university students master six core subjects via personalized AI tutoring, interactive visualizations, and hands-on projects.
- **Target Users:** First-year undergraduates in STEM disciplines.
- **Pedagogical Approach:** Personalized learning paths, contextual prompt suggestions, interactive mind maps, citation-backed responses, embedded multimedia, and project-based learning (PBL).

2. Feature Analysis & Viability

#	Feature	Description	Viability & Implementation
1	Automated Prompt & Context Suggestions	As users type queries (e.g., “Teach me photosynthesis”), real-time AI proposes 2–3 teaching modes (“spoon-feed,” “questioner,” etc.) before generating the answer. Applies equally to uploaded docs/images.	<b>Viability:</b> High. Implementable via a lightweight frontend listener invoking a contextual suggestion service (e.g., OpenAI function calling). The server uses RAG (Retrieval-Augmented Generation) to analyze context and returns suggestion tokens. Suggestion UI layers above the chat input for mode-selection prior to submission.
2	NotebookLM-Style Discover & Mind Maps	“Discover” crawls selected web sources as pseudo-uploads; “Mind Maps” render branching diagrams of topics with clickable nodes revealing text, visuals, and references.	<b>Viability:</b> Proven by Google’s NotebookLM. Interactive mind maps visually summarize sources with branching nodes ( <a href="#">Google Help</a> ). Implementation would use a graph library (e.g., D3.js or Cytoscape.js) plus a backend extractor that ingests URLs

#	Feature	Description	Viability & Implementation
			via Puppeteer and converts content into knowledge graphs.
3	<b>Perplexity-Style Citations</b>	Every AI response includes inline citations and clickable references (websites, PDFs, books).	<b>Viability:</b> High. Perplexity crawls the web at query time, ranks authoritative results, and cites them ( <a href="#">Perplexity AI</a> , <a href="#">TIME</a> ). Evolve can integrate an open-source search API (e.g., Bing's Web Search API) and overlay citation metadata in responses.
4	<b>YouTube API Integration</b>	Query terms trigger YouTube Data API calls; embed top-ranked instructional videos inline.	<b>Viability:</b> Straightforward. Use YouTube Data API v3's <code>search.list</code> endpoint, filter by relevance and educational metadata, then embed via iFrame.
5	<b>Dynamic Diagrams &amp; Visuals</b>	Generate SVG code for conceptual diagrams (e.g., photosynthesis flow, data-structure graphs).	<b>Viability:</b> LLMs (e.g., OpenAI GPT-4) can output Mermaid or raw SVG markup. A client-side SVG renderer (like Mermaid.js) can visualize diagrams on the fly.
6	<b>Web-Sourced Hyperlinks</b>	Supplement responses with direct hyperlinks to topics and resources.	<b>Viability:</b> Combines web search results with citation metadata. Implementation parallels feature #3 but surface links exclusively.
7	<b>Dedicated Programming Pages</b>	Separate "DSA" and "OOP" pages with integrated code editor, compiler, visualizations, and AI code assistant.	<b>Viability:</b> High. Leverage Monaco Editor for code editing, serverless functions (e.g., AWS Lambda / Supabase Edge) for compilation, and custom visualizers (e.g., p5.js) for algorithm animations. AI integration via LangChain agents.

#	Feature	Description	Viability & Implementation
8	Image Retrieval	Fetch relevant images (from Google Images or Unsplash API) for visual learning.	<b>Viability:</b> Use Unsplash or Pexels API for royalty-free images; fallback to Google Custom Search JSON API for broader coverage.
9	LangChain & LangGraph Integration	Utilize LangChain for LLM orchestration; employ LangGraph for stateful, agentic workflows.	<b>Viability:</b> LangChain provides high-level abstractions (chains, tools, agents), while LangGraph offers low-level graph-based orchestration for complex multi-step tasks ( <a href="#">LangChain</a> , <a href="#">IBM</a> ). Recommended: prototype with LangChain; scale with LangGraph where agent state management and debugging are critical.
10	Voice I/O	Multimodal input: speech-to-text prompt entry; text-to-speech response playback.	<b>Viability:</b> Integrate browser Web Speech API for STT/ TTS or use cloud services (e.g., Google Cloud Speech-to-Text, Amazon Polly). Minimal latency for short prompts.
11	Infinite Canvas with Node-Link Notes	Graph-style note linking (akin to Obsidian); AI can summarize interconnected nodes.	<b>Viability:</b> Implement using a canvas library (e.g., Fabric.js) or graph DB (e.g., Neo4j) with a front-end like mermaid-diagrams or React Flow. AI summarization via a LangChain agent that ingests selected node content.
12	Teaching Modes (Tutor, Study Buddy, Questioner, Spoon-feeding, Practical Learning)	Mode selection tailors response style (e.g., Socratic questioning vs. direct explanation vs. project suggestions).	<b>Viability:</b> Simple prompt templates behind the scenes. E.g., system prompt “You are a spoon-feeding tutor: break down ...” vs. “You are a Socratic questioner: ask probing questions ....”
13	Project-Based Learning (PBL) Suggestions	Generate three context-relevant project	<b>Viability:</b> Use prompt engineering to ask LLM for three practical project outlines given topic context. Store

#	Feature	Description	Viability & Implementation
		ideas; support bookmarking selections in user profile (DB). UI for and progress tracking.	bookmarking and displaying progress.
14	Community & Social Features	Optional profiles, friend requests, leaderboards, universal feedback stream moderated by admins.	<b>Viability:</b> Core social graph features (users, profiles, connections) implementable via Supabase Auth + Postgres. Leaderboards via materialized views ranking completed exercises. Universal chat channels with moderation roles.
15	Teacher Dashboard	Subject-specific teacher consoles to upload materials, assign projects, and post announcements.	<b>Viability:</b> Role-based access control (RBAC) via Supabase. Admin UI built with React and shadcn/ui components. File uploads to Supabase Storage.

### 3. Technology Stack Evaluation

Layer	Chosen Tech	Rationale & Alternatives
Frontend	Next.js, React, TypeScript, Tailwind CSS, shadcn/ui	<b>Pros:</b> Server-side rendering, fast HMR, strong TS support. <b>Alternatives:</b> Gatsby (static), Remix (full SSR)
State Management	React Query, Zustand	<b>Pros:</b> Caching, minimal boilerplate. <b>Alt:</b> Redux Toolkit
Backend / API	Node.js (Next.js API Routes) & Python (AI microservices)	<b>Rationale:</b> Node.js for web endpoints; Python for AI workloads (fast model inference with <code>transformers</code> / <code>fastapi</code> ). <b>Alt:</b> All-Python (Django), or Golang microservices
Database	Supabase (Postgres)	<b>Pros:</b> Auth, Realtime, Storage, Row-Level Security out-of-the-box. <b>Alt:</b> Firebase + Cloud Firestore, AWS Amplify

Layer	Chosen Tech	Rationale & Alternatives
<b>AI Orchestration</b>	LangChain + LangGraph	<b>Rationale:</b> Rapid prototyping (LangChain), scalable agents (LangGraph).
<b>Search &amp; Citations</b>	Bing Web Search API / custom crawler	<b>Alt:</b> Algolia with web indexing, ElasticSearch crawler
<b>Video &amp; Images</b>	YouTube Data API, Unsplash API	<b>Alt:</b> Vimeo API, Google Custom Search
<b>Voice I/O</b>	Web Speech API + Cloud TTS/STT	<b>Alt:</b> Azure Cognitive Services
<b>Diagram Rendering</b>	Mermaid.js / custom SVG viewer	<b>Alt:</b> JointJS, D3.js
<b>Auth &amp; Profiles</b>	Supabase Auth	<b>Alt:</b> Auth0, Clerk
<b>Deployment</b>	Vercel (frontend), Fly.io / Railway (backend)	<b>Alt:</b> AWS (Amplify + Lambda), GCP Cloud Run

---

## 4. Recommended Development Methodology

### 1. Modular Architecture:

- **Componentize UI:** Break pages into reusable React components (e.g., <ChatWindow>, <MindMapCanvas>, <CodePlayground>).
- **Microservices:** Separate AI inference (Python) from web API (Node.js) with clear REST/gRPC boundaries.

### 2. Iterative MVPs:

- **Phase 1:** Core chat interface + prompt suggestions + citation integration.
- **Phase 2:** Media embeds (YouTube, images) + diagram generation.
- **Phase 3:** Mind maps + PBL suggestions + teacher dashboard.
- **Phase 4:** Voice I/O + community features + advanced agent workflows (LangGraph).

### 3. Version Control & CI/CD:

- Monorepo managed with TurboRepo or Nx for frontend/backend.
- Automated tests: Jest (unit), Cypress (E2E), Lighthouse (performance).
- CI/CD pipelines on GitHub Actions for test, lint, deploy to staging & production.

- 4. **UX & Accessibility:**
  - o Design system with shaden/ui for consistency.
  - o ARIA roles for screen readers. Voice I/O complement keyboard navigation.
- 5. **Data Privacy & Security:**
  - o Row-Level Security policies in Supabase.
  - o OAuth scopes limited for YouTube/Unsplash.
  - o GDPR and COPPA compliance (no under-13 capture).

---

## 5. Next Steps & Timeline

Sprint	Deliverables	Duration
1	Chat UI, prompt suggestions, basic citation layer	2 weeks
2	Media embeds (videos/images), diagram generation	2 weeks
3	Mind map canvas, PBL module	3 weeks
4	Teacher dashboard, profiles, community feed	3 weeks
5	Voice I/O, LangGraph agent integration	2 weeks
<b>Total</b>	<b>End-to-end MVP</b>	<b>12 weeks</b>

---

## 6. References

- 1. Google NotebookLM Mind Maps feature documentation ([Google Help](#))
  - 2. Perplexity AI real-time search & citation methodology ([Perplexity AI](#), [TIME](#))
  - 3. LangChain & LangGraph overview ([LangChain](#), [IBM](#))
- 

*Prepared by: [Your Name], [Date]*